

IMEN266 Operations Research II

Concept & Proof Companion — Chapter 4 (Markov Chains)

[GAP-FILL version — fill the blanks yourself first]

Fall 2026

How to use this document

This companion *replaces your class notes*: every definition and theorem from the Ch.4 slides is written out here in full, so you can watch and think in class instead of transcribing. It does **not** replace *doing*: for each core proof, (1) work through the GAP-FILL version with pen and paper, (2) check against the FULL version, (3) answer the Checkpoints, and (4) if stuck — or done — use an AI study card. Exams are closed-book; these eight proofs and their checkpoints are exactly the kind of reasoning the exams reward. Items marked **(midterm-style)** mirror questions from past midterms: the subsampling proof (Proof 8), a path-factorisation true/false (Proof 1), P^6 from a given P^3 (Proof 2), a long-run cost (Proof 6), and PageRank with a dead end (Proof 7).

1 Concept digest

1.1 Stochastic processes

A **stochastic process** is an indexed family of random variables $\{X_t, t \in \mathcal{I}\}$: X_t is a random variable for each fixed t , and $\mathcal{I}$ is the **index set** (usually time). Both the index set and the state space can be discrete or continuous:

	index discrete	index continuous
state discrete	Ch. 4: DTMC (inventory at day start)	Ch. 5–7: Poisson process, CTMC, renewal
state continuous	time series (hourly temperature)	Brownian motion (financial engineering)

(verified in notebook: CP3 warm-up)

1.2 Markov chains

$\{X_n, n \in \mathbb{Z}_+\}$ with countable state space $\mathcal{S}$ is a **(discrete-time) Markov chain** if, for all n and all states,

$$P(X_{n+1} = j \mid X_n = i, X_{n-1} = i_{n-1}, \dots, X_0 = i_0) = P(X_{n+1} = j \mid X_n = i) \quad (\text{Markov property})$$

and, in this course, the right-hand side does not depend on n (**time-homogeneity**): $P(X_{n+1} = j \mid X_n = i) = P(X_1 = j \mid X_0 = i) =: P_{ij}$. The **transition probability matrix** $P = (P_{ij})$ is *row-stochastic*: $P_{ij} \geq 0$ and $\sum_j P_{ij} = 1$ for every i . The **transition diagram** draws states as nodes and $P_{ij} > 0$ as arrows.

Modeling recipe (slide p.7/52; every CP3 problem): (1) define X_n , the state at the n -th observation; (2) write $\mathcal{S}$; (3) verify Markov and time-homogeneity (“regardless of the history”, “independent of the past”); (4) derive P — and check that rows sum to 1; (5) draw the diagram. If the natural description has memory (age of a tool, day of a spell), *enlarge the state* until it is Markov (CP3 #6–#7). (verified in notebook: CP3 #1–#7)

1.3 Transient analysis

Initial distribution $a = [P(X_0 = 0), P(X_0 = 1), \dots]$; $p_n = [P(X_n = 0), P(X_n = 1), \dots]$ (row vectors). Then

$$p_1 = aP, \quad p_2 = p_1P = aP^2, \quad \dots, \quad p_n = aP^n \quad (\text{Proof 2}).$$

The n -step transition probabilities $P_{ij}^{(n)} = P(X_n = j \mid X_0 = i)$ satisfy the **Chapman–Kolmogorov equations** $P_{ij}^{(m+n)} = \sum_k P_{ik}^{(m)} P_{kj}^{(n)}$, i.e. $P^{(n)} = P^n$ (matrix power). Path probabilities multiply along the path (Proof 1): $P(X_0 = i_0, \dots, X_n = i_n) = a_{i_0} P_{i_0 i_1} \cdots P_{i_{n-1} i_n}$. Observing the chain every k -th step gives a DTMC with matrix P^k (Proof 8). (*verified in notebook: CP4 #1–#3*)

1.4 Classification of states (slides pp. 29–33/52)

State j is **accessible** from i ($i \rightarrow j$) if $P_{ij}^{(n)} > 0$ for some $n \geq 0$; i and j **communicate** ($i \leftrightarrow j$) if each is accessible from the other. Communication is an equivalence relation (Proof 3), so $\mathcal{S}$ splits into **communicating classes**; the chain is **irreducible** if there is only one class.

Let $f_i = P(\text{ever return to } i \mid X_0 = i)$. State i is **recurrent** if $f_i = 1$ and **transient** if $f_i < 1$; equivalently (Proof 3) i is recurrent iff $\sum_{n \geq 1} P_{ii}^{(n)} = \infty$. Recurrence and transience are **class properties**. A recurrent state is **positive recurrent** if the mean return time $m_i = E[\text{return time} \mid X_0 = i]$ is finite (**null recurrent** otherwise); in a *finite* chain every recurrent state is positive recurrent, and a class is recurrent iff it is *closed* (no arrow leaves it). A state with $P_{ii} = 1$ is **absorbing**.

The **period** of i is $d(i) = \gcd\{n \geq 1 : P_{ii}^{(n)} > 0\}$; i is **aperiodic** if $d(i) = 1$ (a self-loop $P_{ii} > 0$ suffices). Period is also a class property. **Ergodic** = irreducible + positive recurrent + aperiodic (slide p. 34/52: the “desired” chains). (*verified in notebook: CP5 warm-up, DTMC.summary()*)

1.5 Limiting behavior (slide p. 36/52)

Theorem. If the chain is irreducible and positive recurrent, the **stationary equations**

$$\pi = \pi P, \quad \sum_{i \in \mathcal{S}} \pi_i = 1$$

have a unique solution π ; π_i is the long-run fraction of time spent in i , and $\pi_i = 1/m_i$ (Proof 6). If the chain is also aperiodic (ergodic), then $\lim_{n \rightarrow \infty} P(X_n = i) = \pi_i$ for every initial distribution, i.e. every row of P^n converges to π (Proof 4 for two states).

Reading the theorem. Periodic chains (the 3-cycle) still have a unique π but no limit; reducible chains may have limits that depend on the starting state (slide p. 29/52, CP5 warm-up). One of the balance equations is always redundant — replace it by the normalisation (CP5 #2).

1.6 Cost and reward (slide p. 45/52)

If cost $C(i)$ is incurred every time the chain is observed in state i , the long-run average cost per observation is

$$\lim_{n \rightarrow \infty} \frac{1}{n} \sum_{k=1}^n C(X_k) = \sum_{i \in \mathcal{S}} C(i) \pi_i \quad (\text{Proof 6}).$$

Examples: availability (CP5 #1), wet trips $\varphi = p\pi_0$ and $\tau(r) = c\varphi + dr$ (CP5 #3), mean tool age $\sum_i i\pi_i$ and mean lifetime $1/\pi_0$ (CP5 #4), lunch revenue (CP5 #5).

1.7 PageRank (handout Google.PageRank.MC.pdf, in ch04/slides/)

A random surfer on n web pages follows one of the k outgoing links of the current page with probability $1/k$: a DTMC with matrix P ; the rank of page j is π_j . **Dead ends** (no outgoing links: zero row, mass leaks) and **spider traps** (absorbing pages: π concentrates there) break ergodicity. **Taxation** with rate $1 - \beta$ replaces P by

$$G = \beta P + \frac{1 - \beta}{n} \mathbf{1}\mathbf{1}^\top,$$

a strictly positive stochastic matrix, so $\pi = \pi G$ has a unique solution and power iteration $p_{k+1} = p_k G$ converges at rate $\leq \beta$ (Proof 7). Handout numbers: $\pi = (0.4, 0.4, 0.2)$ for the 3-page web; with the spider trap and 30% tax, $\pi \approx (0.26, 0.19, 0.55)$. (verified in notebook: PageRank practicum)

1.8 Chain zoo (the class-problem models)

model (CP3)	state X_n	$ \mathcal{S} $	ergodic?	stationary distribution
machine (#1)	down/up	2	yes	$\pi_{\text{down}} = \frac{b}{a+b}$, $a = P_{\text{du}}$, $b = P_{\text{ud}}$ (Proof 4)
Pohang weather (#2)	S/C/R	3	yes	solve 3×3 : (0.479, 0.311, 0.210)
Himart (#3)	Monday stock	4	yes	numeric; tail probabilities in column "5"
two routes (#4)	# free routes	3	yes	product of two 2-state chains
umbrellas (#5)	# at current place	$r+1$	yes	$\pi_0 = \frac{1-p}{r+1-p}$, $\pi_i = \frac{1}{r+1-p}$ (Proof 5)
tool age (#6)	age	$K+1$	yes	$\pi_i = \pi_0 \prod_{k=1}^i (1-p_k)$, $1/\pi_0 = \text{mean lifetime}$
weather spells (#7)	(type, day of spell)	$2K$	yes	numeric
Kim's lunch (CP5 #5)	restaurant	3	yes	(10, 5, 4)/19

2 Core proofs

Notation: $\mathcal{S}$ finite unless stated; a is the initial distribution; $P_{ij}^{(n)} = P(X_n = j \mid X_0 = i)$. Blanks are the steps *you* must supply in the gap-fill version.

Proof 1 — Path probabilities: the Markov property in action

Theorem. For any states $i_0, \dots, i_n$,

$$P(X_0 = i_0, X_1 = i_1, \dots, X_n = i_n) = a_{i_0} P_{i_0 i_1} P_{i_1 i_2} \cdots P_{i_{n-1} i_n},$$

and consequently $P(X_1 = j) = \sum_i a_i P_{ij}$, i.e. $p_1 = aP$.

Proof. By the multiplication rule (Ch. 1–3, Companion §1.2) applied n times,

$$P(X_0 = i_0, \dots, X_n = i_n) = P(X_0 = i_0) \prod_{k=1}^n P(X_k = i_k \mid \underline{\hspace{10em}}).$$

By the Markov property each conditional probability equals $P(X_k = i_k \mid X_{k-1} = i_{k-1})$, and by time-homogeneity this is $\underline{\hspace{1cm}}$, independent of k . Multiplying gives the product. For the consequence, the events $\{X_0 = i\}$, $i \in \mathcal{S}$, partition Ω , so by the law of total probability $P(X_1 = j) = \sum_i P(X_1 = j \mid X_0 = i)P(X_0 = i) = \sum_i \underline{\hspace{1cm}}$. $\square$

(verified in notebook: CP4 #3(a),(b): $0.3 \cdot 0.25 \cdot 0.3 \cdot 0.1$ matched the simulation)

Checkpoint — answer before moving on.

- In CP4 #3(b), $P[X_3 = 3, X_2 = 2, X_1 = 1]$ has no X_0 : which step of the proof is replaced by a sum, and over what?
- Write the same product for a chain that is Markov but *not* time-homogeneous (transition matrices $P^{(k)}$ at step k).
- Why does the factorisation fail for the weather-spell process if the state is only “sunny/rainy” (CP3 #7)?
- **(midterm-style)** True or false, with a reason: $P(X_6 = 2, X_3 = 1, X_1 = 0 \mid X_0 = 2) = P(X_6 = 2 \mid X_0 = 2) P(X_3 = 1 \mid X_0 = 2) P(X_1 = 0 \mid X_0 = 2)$. Write the correct factorisation.

AI study card — paste one of these into your AI assistant:

- Quiz me one step at a time on why $P(X_0=i_0, \dots, X_n=i_n)$ factorises for a Markov chain. Make me say exactly where the Markov property and where time-homogeneity are used.
- Invent a 3-state process that is NOT Markov, show me a path whose probability does not factorise, and make me compute the true probability.

Proof 2 — Chapman–Kolmogorov and $p_n = aP^n$

Theorem. $P_{ij}^{(m+n)} = \sum_{k \in S} P_{ik}^{(m)} P_{kj}^{(n)}$ for all $m, n \geq 0$; hence $P^{(n)} = P^n$ and $p_n = aP^n$.

Proof. Condition on the state at time m (law of total probability):

$$P_{ij}^{(m+n)} = P(X_{m+n} = j \mid X_0 = i) = \sum_k P(X_{m+n} = j \mid X_m = k, X_0 = i) P(______ \mid X_0 = i).$$

By the Markov property the first factor does not depend on X_0 , and by time-homogeneity it equals $P_{ik}^{(m)}$; the second factor is $P_{kj}^{(n)}$. The sum is the (i, j) entry of the matrix product $P^{(m)} P^{(n)}$. Taking $m = 1$ and inducting on n : $P^{(n)} = P P^{(n-1)} = ___$. Finally $P(X_n = j) = \sum_i P(X_0 = i) P_{ij}^{(n)} = (aP^n)_j$. $\square$

(verified in notebook: CP4 #1–#2: aP^n vs. 200,000 simulated paths)

Checkpoint — answer before moving on.

- Explain in words why $P(X_5 = 2 \mid X_1 = 3) = (P^4)_{32}$ (CP4 #3(c)) — which two properties are used?
- Compute P^2 for the Pohang weather chain by hand and confirm $(P^2)_{SR} = 0.21$.
- True or false: $P_{ik}^{(m)} P_{kj}^{(n)}$ is the probability of going $i \rightarrow k \rightarrow j$ in $m + n$ steps *via* k at time m . What does summing over k do?

- **(midterm-style)** For $P = \begin{pmatrix} 0.5 & 0 & 0.5 \\ 0 & 0 & 1 \\ 0.3 & 0.2 & 0.5 \end{pmatrix}$ Python reports $P^3 = \begin{pmatrix} 0.35 & 0.10 & 0.55 \\ 0.30 & 0.10 & 0.60 \\ 0.33 & 0.12 & 0.55 \end{pmatrix}$. Find

$P(X_{10008} = 1 \mid X_{10002} = 2)$ *without* computing P^6 from scratch. (Which C–K product? Compare with π_1 from Proof 6’s checkpoint.)

AI study card — paste one of these into your AI assistant:

- Play a skeptical student: I will prove the Chapman-Kolmogorov equations; interrupt whenever I use a property without naming it.
- Give me a 3x3 transition matrix and ask me for $P(X_4=j \mid X_1=i)$ two ways: by matrix power and by summing over paths. Check that my two answers agree.

Proof 3 — Recurrence: the $\sum_n P_{ii}^{(n)}$ criterion and class property

Theorem. (a) State i is recurrent iff $\sum_{n \geq 1} P_{ii}^{(n)} = \infty$. (b) If i is recurrent and $i \leftrightarrow j$, then j is recurrent. (c) A finite irreducible chain is recurrent.

Proof of (a). Let $N = \sum_{n \geq 1} \mathbf{1}\{X_n = i\}$ be the number of returns to i . Each time the chain is in i it returns again with probability f_i , independently of the past (Markov property), so $P(N \geq k \mid X_0 = i) = ______$ and by the tail-sum formula (Ch. 1–3, Proof 3)

$$E[N \mid X_0 = i] = \sum_{k \geq 1} f_i^k = \frac{f_i}{1 - f_i} \quad (f_i < 1), \quad E[N \mid X_0 = i] = \infty \quad (f_i = 1).$$

On the other hand, by linearity of expectation $E[N \mid X_0 = i] = \sum_{n \geq 1} E[\mathbf{1}\{X_n = i\} \mid X_0 = i] = \sum_{n \geq 1} ______$. So the sum is infinite iff $f_i = 1$. $\square$

Proof of (b). Since $i \leftrightarrow j$ there are k, m with $P_{ij}^{(k)} > 0$ and $P_{ji}^{(m)} > 0$. One way to go $j \rightarrow j$ in $m + n + k$ steps is $j \rightarrow i$ in m , $i \rightarrow i$ in n , $i \rightarrow j$ in k , so

$$P_{jj}^{(m+n+k)} \geq P_{ji}^{(m)} ______ P_{ij}^{(k)}.$$

Summing over n : $\sum_n P_{jj}^{(m+n+k)} \geq P_{ji}^{(m)} P_{ij}^{(k)} \sum_n P_{ii}^{(n)} = ______ by (a), so j is recurrent by (a). $\square$$

Proof of (c). If all states were transient, each is visited only finitely often, so after some finite time the chain would be in *no* state — impossible. Hence some state is recurrent, and by (b) and irreducibility $______ states are. $\square$$

Remark. In a finite chain a class is recurrent iff it is closed; transient classes are those with an arrow leaving them (slide p. 29/52: $\{3, 5\}$ leaks into the absorbing state 4). (verified in notebook: CP5 warm-up, summary())

Checkpoint — answer before moving on.

- Where did the proof of (a) use the Markov property? (Why are successive returns independent trials with the same success probability?)
- Symmetric random walk on $\mathbb{Z}$ is recurrent but its mean return time is infinite (null recurrent). Why can this not happen in a finite chain?
- Classify every state of the two matrices on slide p. 29/52 and give the period of each class.

AI study card — paste one of these into your AI assistant:

- Ask me to prove that recurrence is a class property. Do not accept "obvious"; make me write the inequality between n -step probabilities.
- Give me five small transition matrices; for each I must state the classes, which are recurrent, and the period. Grade me and explain any mistake with a path argument.

Proof 4 — The two-state chain: P^n in closed form and its limit

Theorem. Let $P = \begin{pmatrix} 1-a & a \\ b & 1-b \end{pmatrix}$ with $0 < a+b < 2$ (states 0, 1; $a = P_{01}$, $b = P_{10}$). Then with $\lambda = 1 - a - b$,

$$P_{00}^{(n)} = \frac{b}{a+b} + \frac{a}{a+b} \lambda^n, \quad P_{10}^{(n)} = \frac{b}{a+b} - \frac{b}{a+b} \lambda^n, \quad \text{so } P^n \rightarrow \begin{pmatrix} \frac{b}{a+b} & \frac{a}{a+b} \\ \frac{b}{a+b} & \frac{a}{a+b} \end{pmatrix} = \begin{pmatrix} \pi \\ \pi \end{pmatrix}.$$

Proof. Let $x_n = P_{00}^{(n)}$. By Chapman–Kolmogorov (condition on the state at time n),

$$x_{n+1} = P_{00}^{(n)} P_{00} + P_{01}^{(n)} P_{10} = x_n(1-a) + \frac{b}{a+b} = b + \lambda x_n.$$

The fixed point of $x \mapsto b + \lambda x$ is $x^* = \frac{b}{1-\lambda} = \frac{b}{a+b}$, and subtracting $x^* = b + \lambda x^*$ gives $x_{n+1} - x^* = \lambda(x_n - x^*)$, hence $x_n - x^* = \lambda^n(x_0 - x^*)$ with $x_0 = P_{00}^{(0)} = 1$. So $x_n = \frac{b}{a+b} + \lambda^n(1 - \frac{b}{a+b}) = \frac{b}{a+b} + \frac{a}{a+b} \lambda^n$. The same recursion with $x_0 = P_{10}^{(0)} = 0$ gives the second formula. Since $0 < a+b < 2$ means $|\lambda| < 1$, $\lambda^n \rightarrow 0$ and both rows converge to $\pi = (\frac{b}{a+b}, \frac{a}{a+b})$, which indeed solves $\pi = \pi P$. $\square$

Machine (CP5 #1). down= 0, up= 1: $a = 0.97$, $b = 0.02$, $\pi_{\text{down}} = \frac{0.02}{0.99} \approx 0.0202$, $\lambda = 0.01$ — the chain forgets its start almost immediately. (verified in notebook: CP5 #1 experiment: speed of forgetting = $|\lambda|$)

Checkpoint — answer before moving on.

- What happens when $a+b=1$? When $a+b=2$? Interpret both with the transition diagram and with λ .
- λ is the second eigenvalue of P (the first is 1). Verify by computing the characteristic polynomial.
- Write the cut-balance equation $\pi_0 a = \pi_1 b$ and derive π from it in one line — then compare with the limit above.

AI study card — paste one of these into your AI assistant:

- Walk me through the two-state Markov chain P^n derivation by asking one question per step; I must supply the recursion and the fixed point myself.
- Give me three two-state chains ($a+b<1$, $a+b=1$, $a+b>1$) and make me predict, before computing, whether P^n converges monotonically, instantly, or with oscillation.

Proof 5 — Solving the stationary equations: the umbrella chain

Theorem. For the umbrella chain with $r \geq 1$ umbrellas and rain probability $p \in (0, 1)$ (CP3 #5), the unique stationary distribution is

$$\pi_0 = \frac{1-p}{r+1-p}, \quad \pi_1 = \cdots = \pi_r = \frac{1}{r+1-p},$$

and the long-run fraction of wet trips is $\varphi = p \pi_0 = \frac{p(1-p)}{r+1-p}$.

Proof. Transitions: $0 \rightarrow r$ w.p. 1; for $i \geq 1$, $i \rightarrow r - i$ w.p. $1 - p$ and $i \rightarrow r - i + 1$ w.p. p . Write the balance equation $\pi_j = \sum_i \pi_i P_{ij}$ for each *target* j by asking “who leads to j ?”:

$$\begin{aligned} j = r : \quad \pi_r &= \pi_0 \cdot 1 + \pi_1 \cdot _ \quad (\text{from } 0, \text{ and from } 1 \text{ in rain}), \\ 1 \leq j \leq r - 1 : \quad \pi_j &= \pi_{r-j} (1 - p) + \pi_{r-j+1} p, \\ j = 0 : \quad \pi_0 &= \pi_r \cdot _ \quad (\text{from } r \text{ with no rain}). \end{aligned}$$

Try $\pi_1 = \dots = \pi_r = c$. The middle equations become $c = c(1 - p) + cp$, true. The last gives $\pi_0 = _$, and then the first reads $c = c(1 - p) + cp$ — true again. Normalising, $c(1 - p) + rc = 1 \Rightarrow c = _$. The chain is finite and irreducible (every state leads to r , and r leads to every state via $\overline{r - i}$ and $r - i + 1$), so this solution is the unique one. Finally she gets wet on a trip iff it rains *and* she is in state 0; rain is independent of the past, so $\varphi = P(\text{rain}) P(X_n = 0) \rightarrow p \pi_0$. $\square$

Cost (CP5 #3). $\tau(r) = c\varphi(r) + dr$ is convex in r ; the continuous minimiser is $r^* = \sqrt{cp(1 - p)/d} - (1 - p)$, and the optimal integer is one of its neighbours.

Checkpoint — answer before moving on.

- Which single balance equation was redundant here? Show that adding all $r + 1$ equations gives $0 = 0$.
- Why is φ maximal at $p = \frac{1}{2}$ and not at $p \rightarrow 1$? Give the one-sentence reason from the formula and from the story.
- Verify aperiodicity: find the self-loop for even r and for odd r .

AI study card — paste one of these into your AI assistant:

- Ask me to set up (not solve) the balance equations for the umbrella chain with $r=3$; check each equation against “who leads into this state”.
- Claim (falsely) that the fraction of wet trips is π_0 . Make me find the missing factor and explain why the two events are independent.

Proof 6 — Long-run fractions, mean return times, and the cost formula

Theorem. Let the chain be irreducible and positive recurrent on a finite $\mathcal{S}$, with m_i the mean return time to i . Then, with probability 1,

$$(a) \quad \frac{1}{n} \sum_{k=1}^n \mathbf{1}\{X_k = i\} \rightarrow \frac{1}{m_i} = \pi_i, \quad (b) \quad \frac{1}{n} \sum_{k=1}^n C(X_k) \rightarrow \sum_{i \in \mathcal{S}} C(i) \pi_i.$$

Proof of (a) (time fraction = $1/m_i$). Let $T_1 < T_2 < \dots$ be the successive visit times to i . By the Markov property and homogeneity, the gaps $T_{\ell+1} - T_\ell$ are i.i.d. with mean $_$ (each restarts the chain from i). Let $N(n)$ be the number of visits by time n ; then $T_{N(n)} \leq n < T_{N(n)+1}$, so

$$\frac{T_{N(n)}}{N(n)} \leq \frac{n}{N(n)} < \frac{T_{N(n)+1}}{N(n)+1} \cdot \frac{N(n)+1}{N(n)}.$$

Since i is recurrent, $N(n) \rightarrow \infty$, and by the strong law of large numbers both outer terms converge to $_$. Hence $n/N(n) \rightarrow m_i$, i.e. $N(n)/n \rightarrow 1/m_i$. That this limit is the solution π_i of $\pi = \pi P$ is the content of the limiting theorem (§1.5), which we accept. $\square$

Proof of (b). Group the terms by state:

$$\frac{1}{n} \sum_{k=1}^n C(X_k) = \sum_{i \in \mathcal{S}} C(i) \frac{1}{n} \sum_{k=1}^n \text{—————} \rightarrow \sum_{i \in \mathcal{S}} C(i) \pi_i,$$

using (a) for each of the *finitely many* states (the limit of a finite sum is the sum of the limits). $\square$

Consequences. Mean return time $m_i = 1/\pi_i$ (CP5 #1: a down period recurs every $1/\pi_{\text{down}} \approx 49.5$ days; CP5 #4: mean tool lifetime = $1/\pi_0$). Time averages (Kim's own spending) and ensemble averages (revenue per student) coincide — *ergodicity* (CP5 #5).

Checkpoint — answer before moving on.

- In (a), where exactly is positive recurrence (finite m_i) needed for the sandwich argument to give a finite, nonzero limit?
- CP5 #4: with X the lifetime of a tool, the mean lifetime is $\sum_{k \geq 1} k P(X = k)$ (slide p. 48/52), $\sum_{k \geq 0} P(X > k)$ (tail sum) and $1/\pi_0$. Explain why all three are the same number without computing any of them.
- Part (b) exchanged a limit with a finite sum. Where could it fail for an infinite state space with unbounded C ?
- **(midterm-style)** With the 3-state P of Proof 2's last checkpoint, let X_n be the number of movies rented on the weekly trip n , at 1,500 won per movie plus a 500 won bus fare. Find π and the long-run average weekly cost. (Answer: $\pi = (\frac{1}{3}, \frac{1}{9}, \frac{5}{9})$, cost $\approx 2,333$ won.)

AI study card — paste one of these into your AI assistant:

- Quiz me on why the long-run fraction of time in state i equals $1/m_i$. Force me to write the sandwich inequality with the visit times and to say where the SLLN is used.
- Give me a 3-state chain with costs; I compute the long-run average cost by hand and you check it. Then ask me whether the answer changes with the initial state, and why.

Proof 7 — PageRank: taxation makes the surfer ergodic (exam topic)

Theorem. Let P be an $n \times n$ row-stochastic matrix (dead ends replaced by uniform rows) and $0 < \beta < 1$. Then $G = \beta P + \frac{1-\beta}{n} \mathbf{1}\mathbf{1}^\top$ is row-stochastic and strictly positive, hence ergodic, so $\pi = \pi G$, $\pi \mathbf{1} = 1$ has a unique solution. Moreover, for any probability row vector p_0 and $p_{k+1} = p_k G$,

$$\|p_{k+1} - \pi\|_1 \leq \beta \|p_k - \pi\|_1, \quad \text{so} \quad \|p_k - \pi\|_1 \leq \beta^k \|p_0 - \pi\|_1 \rightarrow 0.$$

Proof. Row sums: $\sum_j G_{ij} = \beta \cdot 1 + (1 - \beta) \cdot \frac{n}{n} = 1$, and every entry is at least $\frac{1-\beta}{n} > 0$. A positive matrix makes every state accessible from every other in one step (irreducible) and gives every state a self-loop (_____); a finite irreducible chain is positive recurrent (Proof 3(c)). So the chain is ergodic and π is unique.

For the rate, let $e_k = p_k - \pi$. Both p_k and π sum to 1, so $e_k \mathbf{1} = 0$. Then, using $\pi = \pi G$,

$$e_{k+1} = p_k G - \pi G = e_k G = \beta e_k P + \frac{1-\beta}{n} (e_k \mathbf{1}) \mathbf{1}^\top = \beta e_k P.$$

Finally $\|e_k P\|_1 = \sum_j |\sum_i (e_k)_i P_{ij}| \leq \sum_i |(e_k)_i| \sum_j P_{ij} = \text{_____}$ because P is row-stochastic. Hence $\|e_{k+1}\|_1 \leq \beta \|e_k\|_1$. $\square$

Reading it. The web's structure never entered the rate: with $\beta = 0.85$, about 50 iterations give 10^{-4} accuracy on *any* web (practicum §5). A page with no incoming links has exactly $\pi_j = (1 - \beta)/n$ (practicum §6). Without taxation ($\beta = 1$) the bound is useless, and spider traps/dead ends break uniqueness (practicum §3). (verified in notebook: PageRank practicum: $\pi = (0.4, 0.4, 0.2)$; with tax $(0.26, 0.19, 0.55)$; $|\lambda_2| \leq \beta$)

Checkpoint — answer before moving on.

- Prove $\pi_j = (1 - \beta)/n$ for a page with no incoming links directly from $\pi = \pi G$.
- Why does a dead end have to be repaired *before* taxation? What goes wrong with the row sums otherwise?
- The 3-cycle (period 3) has no limit; after taxation it does. Which line of the proof kills the period?
- **(midterm-style)** Sites A, B, C: A links to A, B, C; B has no outbound links; C links to A and B. Assign the transition probabilities (repair the dead end), tax with $\beta = 0.85$, and rank. Two sites tie — say which and why before computing (practicum §6, Drill B).

AI study card — paste one of these into your AI assistant:

- Quiz me step by step: show that the taxed PageRank matrix is ergodic and that power iteration contracts the L1 error by the factor beta. Make me justify `e_k 1 = 0`.
- Give me a 4-page web with one spider trap. I must (i) compute pi without tax, (ii) explain what is wrong, (iii) compute pi with beta=0.8 by hand for one iteration. Check every number.

Proof 8 — Subsampling a Markov chain (midterm-style)

Theorem. Let $\{X_n\}$ be a time-homogeneous DTMC with matrix P and fix $k \geq 1$. Then $Y_n = X_{kn}$ ($n \geq 0$) is a time-homogeneous DTMC with transition matrix P^k . In particular $X_0, X_2, X_4, \dots$ is a Markov chain.

Proof. Step 1 (the Markov property reaches r steps ahead and ignores any past event). Let A be any event determined by $X_0, \dots, X_{m-1}$ with $P(X_m = i, A) > 0$. Summing over the intermediate path $X_{m+1} = j_1, \dots, X_{m+r-1} = j_{r-1}$ and applying the one-step Markov property together with time-homogeneity at each step,

$$P(X_{m+r} = j \mid X_m = i, A) = \sum_{j_1, \dots, j_{r-1}} P_{ij_1} P_{j_1 j_2} \cdots P_{j_{r-1} j} = \text{_____} = (P^r)_{ij},$$

which does not depend on A (nor on m). *Step 2.* Take $m = kn$, $r = k$ and $A = \{X_{k(n-1)} = i_{n-1}, \dots, X_0 = i_0\}$, which is determined by $X_0, \dots, X_{kn-1}$:

$$P(Y_{n+1} = j \mid Y_n = i, Y_{n-1} = i_{n-1}, \dots, Y_0 = i_0) = P(X_{k(n+1)} = j \mid X_{kn} = i, A) = \text{_____}.$$

The right-hand side depends only on (i, j) : this is the Markov property for $\{Y_n\}$, and the transition matrix is the same for every n , so $\{Y_n\}$ is time-homogeneous with matrix _____ (a stochastic matrix, since products of stochastic matrices are stochastic). $\square$

Checkpoint — answer before moving on.

- Where exactly was time-homogeneity used, and what would change without it (write the matrix of $\{Y_n\}$ at step n)?
- Is $X_0, X_1, X_3, X_6, X_{10}, \dots$ (gaps 1, 2, 3, 4, ...) a Markov chain? Is it time-homogeneous?
- Subsample the 3-cycle (period 3) with $k = 3$: what is P^3 , and what does the result say about periodic chains and subsampling? Is $\{Y_n\}$ still irreducible?

AI study card — paste one of these into your AI assistant:

- Quiz me: prove that $X_0, X_2, X_4, \dots$ is a Markov chain. Make me define "the past" of the subsampled chain as an event about the original chain and check that the Markov property handles it.
- Give me a 3-state P and ask me for the transition matrix of the chain observed every second step. Then ask whether its stationary distribution changes. Check my answers.

Where to next. CP5 wrap-up → these proofs (gap-fill first) → Checkpoint quiz at the start of the next lecture → HW 2 Part A (closed-book style, no AI on the first attempt) → PageRank practicum §6 drill.